

# Better Online Learning

Su Chen

September 18, 2016

## Part A

```
Sto_Gradient_backtracking = function(x, y, m, alpha0, rho, c, lambda, iter_max) {
  N = nrow(x)
  P = ncol(x)
  # initialize all the variables
  beta = matrix(0, P, iter_max)
  nll_ini = Neg_ll(rep(0, P), x, y, m)
  nll = rep(nll_ini, iter_max)
  nll_avg = rep(nll_ini, iter_max)
  nll_ex_avg = rep(nll_ini, iter_max)
  iter = 1

  while (iter < iter_max) {
    index = sample(1:N, 1) #sample 1 data point from the whole data set
    xi = x[index, , drop = F]
    yi = y[index]
    mi = m[index, ]
    g = Gradient_Cal(beta[, iter], xi, yi, mi)
    desc_direction = -g
    # desc_direction = desc_direction / sqrt(sum(desc_direction^2))
    step = Backtracking(f = Neg_ll, df = g, t = beta[, iter], p = desc_direction,
      alpha0, rho, c, xi, yi, mi)
    beta[, iter + 1] = beta[, iter] + step * desc_direction #update beta
    nll[iter + 1] = Neg_ll(beta[, iter + 1], x, y, m) #calculate neg log likelihood for whole data
    nll_xi = Neg_ll(beta[, iter + 1], xi, yi, mi) * N #calculate neg log likelihood for single data
    nll_avg[iter + 1] = (nll_xi + iter * nll_avg[iter]) / (iter + 1)
    nll_ex_avg[iter + 1] = lambda * nll_xi + (1 - lambda) * nll_ex_avg[iter]
    iter = iter + 1 #keep track of iteration
  }
  return(list(iter, beta[, 1:iter], nll[1:iter], nll_avg[1:iter], nll_ex_avg[1:iter]))
}
```

## Part B

```
Ada_Grad = function(x, y, m, step, lambda, iter_max) {
  N = nrow(x)
  P = ncol(x)
  # initialize all the variables
  beta = matrix(0, P, iter_max)
  nll_ini = Neg_ll(rep(0, P), x, y, m)
  nll = rep(nll_ini, iter_max)
  nll_avg = rep(nll_ini, iter_max)
  nll_ex_avg = rep(nll_ini, iter_max)
  iter = 1
  G = rep(0, P)
```

```

while (iter < iter_max) {
  index = sample(1:N, 1) #sample 1 data point from the whole data set
  xi = x[index, , drop = F]
  yi = y[index]
  mi = m[index, ]
  g = Gradient_Cal(beta[, iter], xi, yi, mi)
  G = G + g^2
  beta[, iter + 1] = beta[, iter] - step * g/(sqrt(G) + 1e-08) #update beta
  nll[iter + 1] = Neg_ll(beta[, iter + 1], x, y, m) #calculate neg log likelihood for whole data
  nll_xi = Neg_ll(beta[, iter + 1], xi, yi, mi) * N #calculate neg log likelihood for single data
  nll_avg[iter + 1] = (nll_xi + iter * nll_avg[iter])/(iter + 1)
  nll_ex_avg[iter + 1] = lambda * nll_xi + (1 - lambda) * nll_ex_avg[iter]
  iter = iter + 1 #keep track of iteration
}
return(list(iter, beta[, 1:iter], nll[1:iter], nll_avg[1:iter], nll_ex_avg[1:iter]))
}

```

```

N = 500
P = 3
X_sim = as.matrix(cbind(rep(1, N), rnorm(N), rnorm(N)))
beta_sim = c(1, 2, -3)
W_sim = 1/(1 + exp(-X_sim %*% beta_sim))
Y_sim = matrix(rbinom(N, 1, W_sim), N, 1)
M_sim = matrix(1, N, 1)

```

```

# run glm and save output for comparision later
glm_sim = glm(Y_sim ~ X_sim - 1, family = "binomial")
result_glm = as.vector(glm_sim$coefficients)
nll_glm = Neg_ll(result_glm, X_sim, Y_sim, M_sim)

```

```

# stochastic gradient decsent with different step size
Iter_Max = 5000

```

```

Step_constant = rep(0.01, Iter_Max)

```

```

Step_RM = rep(0, Iter_Max)
for (i in 1:Iter_Max) {
  Step_RM[i] = RM_Step_Cal(C = 1, t = i, t0 = 1, alpha = 0.5)
}

```

```

result_sgd_constant = Sto_Gradient_Desc(x = X_sim, y = Y_sim, m = M_sim, step = Step_constant,
  lambda = 0.01, iter_max = Iter_Max)

```

```

result_sgd_rm = Sto_Gradient_Desc(x = X_sim, y = Y_sim, m = M_sim, step = Step_RM,
  lambda = 0.01, iter_max = Iter_Max)

```

```

result_sgd_backtracking = Sto_Gradient_backtracking(x = X_sim, y = Y_sim, m = M_sim,
  alpha0 = 0.1, rho = 0.9, c = 0.01, lambda = 0.01, iter_max = Iter_Max)

```

```

result_ada_grad = Ada_Grad(x = X_sim, y = Y_sim, m = M_sim, step = 1, lambda = 0.01,
  iter_max = Iter_Max)

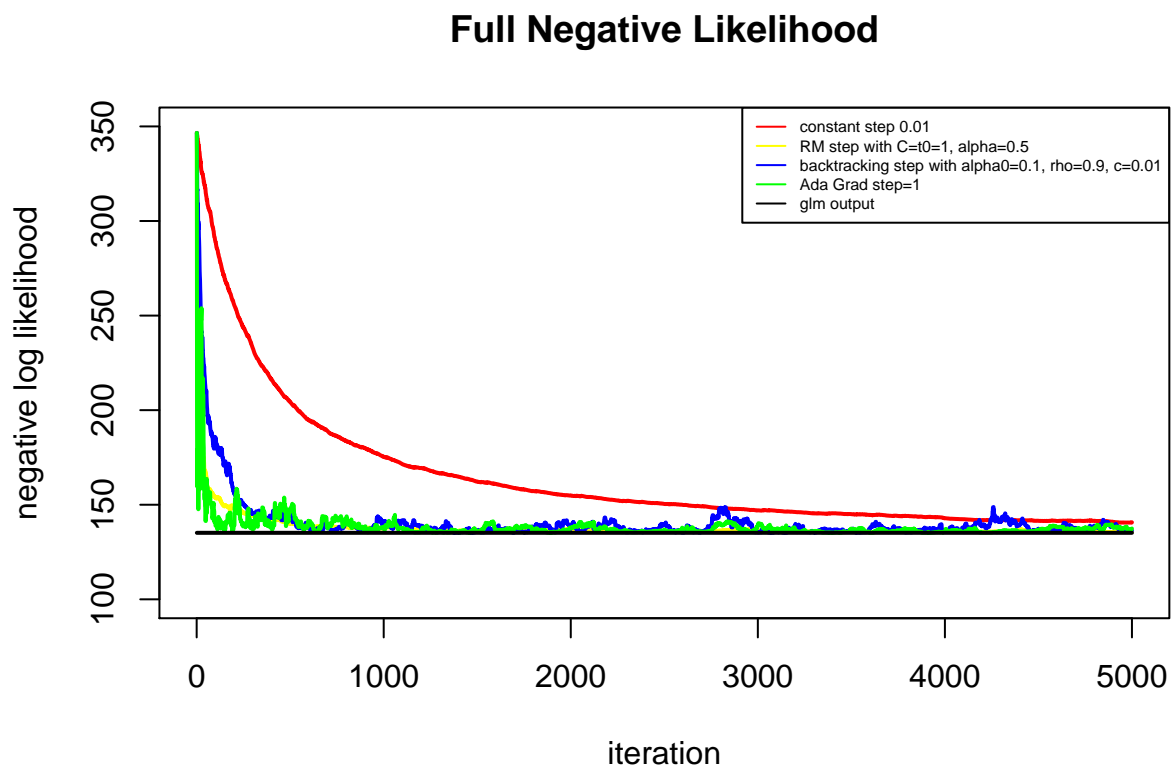
```

```

### plot likelihood to compare convergence for different step size

plot_x = 1:Iter_Max
plot(x = plot_x, y = result_sgd_constant[[3]], type = "l", xlab = "iteration",
     ylab = "negative log likelihood", ylim = c(100, 350), col = "red", lwd = 2,
     main = "Full Negative Likelihood")
lines(x = plot_x, y = result_sgd_rm[[3]], col = "yellow", lwd = 2)
lines(x = plot_x, y = result_sgd_backtracking[[3]], col = "blue", lwd = 2)
lines(x = plot_x, y = result_ada_grad[[3]], col = "green", lwd = 2)
lines(x = c(1, Iter_Max), y = rep(nll_glm, 2), col = "black", lwd = 2)
legend("topright", cex = 0.5, c("constant step 0.01", "RM step with C=t0=1, alpha=0.5",
    "backtracking step with alpha0=0.1, rho=0.9, c=0.01", "Ada Grad step=1",
    "glm output"), col = c("red", "yellow", "blue", "green", "black"), lty = c(1,
    1, 1, 1, 1))

```




```

plot(x = plot_x, y = result_sgd_constant[[4]], type = "l", xlab = "iteration",
     ylab = "negative log likelihood", ylim = c(100, 350), col = "red", lwd = 2,
     main = "Running Avg Negative Likelihood")
lines(x = plot_x, y = result_sgd_rm[[4]], col = "yellow", lwd = 2)
lines(x = plot_x, y = result_sgd_backtracking[[4]], col = "blue", lwd = 2)
lines(x = plot_x, y = result_ada_grad[[4]], col = "green", lwd = 2)
lines(x = c(1, Iter_Max), y = rep(nll_glm, 2), col = "black", lwd = 2)
legend("topright", cex = 0.5, c("constant step 0.01", "RM step with C=t0=1, alpha=0.5",
    "backtracking step with alpha0=0.1, rho=0.9, c=0.01", "Ada Grad step=1",
    "glm output"), col = c("red", "yellow", "blue", "green", "black"), lty = c(1,
    1, 1, 1, 1))

```


1, 1, 1, 1))

## Running Avg Negative Likelihood



```
plot(x = plot_x, y = result_sgd_constant[[5]], type = "l", xlab = "iteration",
     ylab = "negative log likelihood", ylim = c(100, 350), col = "red", lwd = 2,
     main = "Exp Weighted Avg Negative Likelihood")
lines(x = plot_x, y = result_sgd_rm[[5]], col = "yellow", lwd = 2)
lines(x = plot_x, y = result_sgd_backtracking[[5]], col = "blue", lwd = 2)
lines(x = plot_x, y = result_ada_grad[[5]], col = "green", lwd = 2)
lines(x = c(1, Iter_Max), y = rep(nll_glm, 2), col = "black", lwd = 2)
legend("topright", cex = 0.5, c("constant step 0.01", "RM step with C=t0=1, alpha=0.5",
    "backtracking step with alpha0=0.1, rho=0.9, c=0.01", "Ada Grad step=1",
    "glm output"), col = c("red", "yellow", "blue", "green", "black"), lty = c(1,
1, 1, 1, 1))
```

## Exp Weighted Avg Negative Likelihood



### Part C

```
library(Rcpp)
library(RcppEigen)
library(Matrix)
sourceCpp(file = "C:/Users/schen/Dropbox/toChensu/Stats/2016Fall/Big Data/Assignment4/Ada_Grad.cpp")
sourceCpp(file = "C:/Users/schen/Dropbox/toChensu/Stats/2016Fall/Big Data/Assignment4/sgdlogit.cpp")

x = readRDS("C:/Users/schen/Dropbox/toChensu/Stats/2016Fall/Big Data/Assignment4/url_X.rds")
y = readRDS("C:/Users/schen/Dropbox/toChensu/Stats/2016Fall/Big Data/Assignment4/url_y.rds")
x = Matrix(x, sparse = TRUE)
x = t(x)
beta0 = rep(0, nrow(x))

time = proc.time()
result_JS = sparsesgd_logit(x, y, M, eta = 0.01, npass = 1, beta0, lambda = 0.1,
  discount = 0.01)
time_JS = proc.time() - time
plot(result_JS$null_tracker, type = "l", log = "xy", xlab = "iteration", ylab = "Exp Weighted Avg negative log likelihood",
  main = "Stochastic Gradient Decsent JS code")

time = proc.time()
result_SC = Ada_Grad_sparse(x, y, beta0, step = 1e-08, npass = 1, alpha = 0.01,
```

```

    lambda = 0.1)
time_SC = proc.time() - time
plot(result_SC$null_ex_avg, type = "l", log = "xy", xlab = "iteration", ylab = "Exp Weighted Avg negative",
     main = "Stochastic Gradient Decsent SC code")

# compare results of James Scott code with mine code
Output = matrix(0, 6, 2)
dimnames(Output) = list(c("run time", "intercept", "min beta", "median beta",
    "mean beta", "max beta"), c("JS code", "SC code"))
Output[, 1] = c(time_JS[3], result_JS$alpha, min(result_JS$beta), median(result_JS$beta),
    mean(result_JS$beta), max(result_JS$beta))
Output[, 2] = c(time_SC[3], result_SC$intercept, min(result_SC$beta), median(result_SC$beta),
    mean(result_SC$beta), max(result_SC$beta))
Output

```